Run report: Orchestral Input Pipeline: Structured Prompt Conditioning for Full-Ensemble AI Music Generation

Smoke test only. Quick mode checks that the package works. Its numbers are not study results and must not be reported as findings.

Author: KEYUR SHAILESHKUMAR PATEL (ORCID https://orcid.org/0009-0008-5261-043X)

Affiliations: (1) MSc Artificial Intelligence student, Department of Computer Science, Brunel University London, UK; (2) Independent researcher

Conducted independently; not an official Brunel University London project.

Run: 20261028-224832_quick_f9df5d2c (quick mode)

Started 2026-10-28T22:48:32+00:00, finished 2026-10-28T22:48:38+00:00

Config SHA-256: f9df5d2c0a282994825eeb2c7f54774d7cd3e59a41449592b1124ad6f01d4998

Requests SHA-256: 78b651e6688c6520c115aa3448807a3faaf9803ee6078ecc7780dc60f1f6266c

Pieces evaluated: 8 (failed: 0)

Human participants: none (fully computational study)

Every file in this folder is listed with its SHA-256 hash in SHA256SUMS.

Results in plain words

midi: dynamics_accuracy could not be tested (4 pairs; the test needs at least 6).

midi: structure_f could not be tested (4 pairs; the test needs at least 6).

Mean scores by generator and condition

generator	condition	dynamics_accuracy	structure_f	structure_chance	dynamic_range_db	range_violations
midi	pipeline	0.942	0.917	0.719	22.261	0.000
midi	plain	-0.267	0.733	0.719	1.057	58.750

structure_chance is the boundary score expected from random boundaries; structure_f should be read against it.

Confirmatory tests (preregistered)

family	generator	comparison	metric	n_pairs	median_a	median_b	median_diff	ci95_low	ci95_high	W	p	p_holm	significant	rank_biserial
confirmatory	midi	pipeline vs plain	dynamics_accuracy	4	0.965	-0.255	1.154	0.944	1.587	-	-	-	None	1.000
confirmatory	midi	pipeline vs plain	structure_f	4	1.000	0.667	0.167	0.000	0.400	-	-	-	None	1.000

Secondary tests

family	generator	comparison	metric	n_pairs	median_a	median_b	median_diff	ci95_low	ci95_high	W	p	p_holm	significant	rank_biserial
secondary	midi	pipeline vs plain	dynamic_range_db	4	20.498	1.179	19.056	17.401	29.304	-	-	-	None	1.000

Agreement between measures (Spearman)

measure_a	measure_b	n	spearman_rho	p
dynamics_accuracy	structure_f	8	0.548	0.160

Intent-parser accuracy on the benchmark

Mean family F1 1.000, form correct 4/4, dynamic shape correct 4/4 (optimistic: rules were written with these requests).

Held-out requests: none yet (add them to config/heldout_requests.yaml before the main run).

Frechet Audio Distance (CLAP embeddings, lower is closer to real recordings)

No rows.

Limitations

All measures are automatic proxies for musical quality; no human listeners were involved.

The AI judge is itself a model with its own biases; it is reported as a secondary measure only.

CLAP similarities depend on the CLAP model's training data, which under-represents some styles (for example Indian classical music).

Pieces are short (default 30 s), so large-scale form is tested only in compressed form.

The intent parser was developed alongside the main benchmark; its accuracy there is optimistic. Parser accuracy on the held-out requests is the fair estimate.

Closed services (Suno) cannot be re-run exactly and are excluded from confirmatory tests.

Reproduce and verify

python run_all.py verify results/20261028-224832_quick_f9df5d2c